

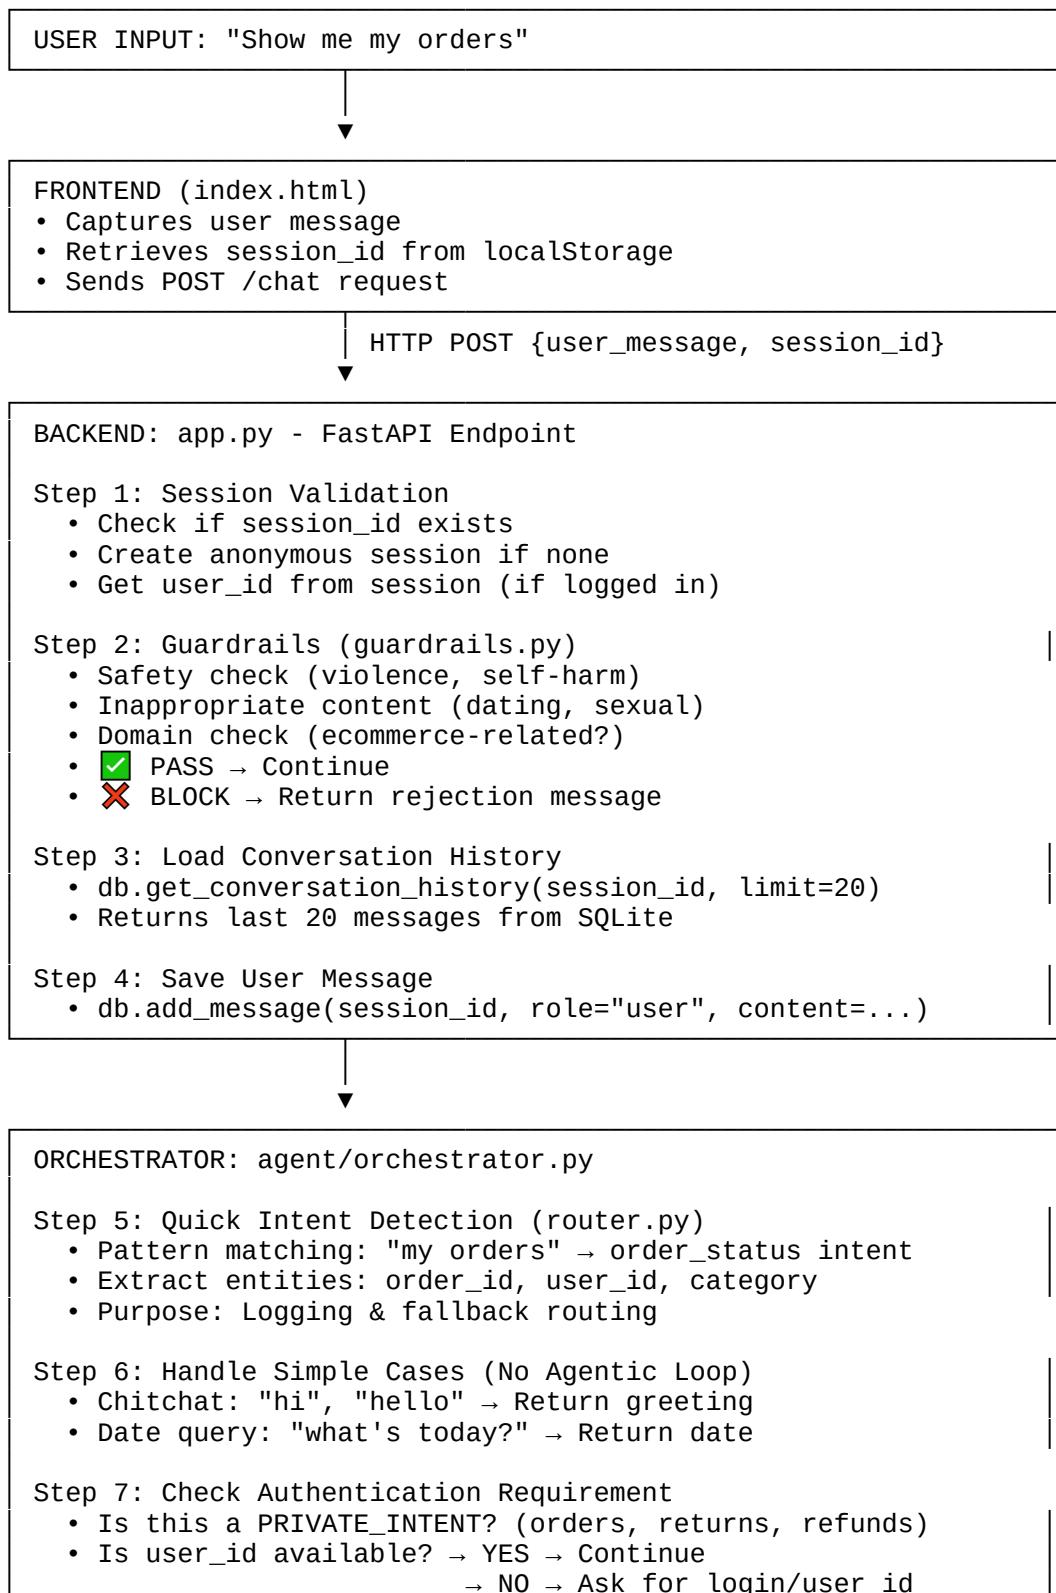
```
# 📚 **Complete System Architecture & Technical Deep Dive**
```

This is your **preparation guide** with all technical details!

```
## ** 1 Complete Workflow: How a Prompt is Treated**
```

```
### **End-to-End Flow Diagram:**
```

```



## Step 8: AGENTIC LOOP (Main Intelligence)

- A. Configure Gemini with Function Calling
  - Load GEMINI\_API\_KEY
  - Create GenerativeModel with tool definitions
- B. Build System Prompt
  - Include user\_id context (if logged in)
  - List available tools
  - Add behavioral guidelines
- C. Format Conversation History
  - Convert DB format → Gemini format
  - user messages → {"role": "user", ...}
  - assistant messages → {"role": "model", ...}
- D. Start Agentic Loop (Max 10 iterations)
  - Iteration 1:
    - Send: system\_prompt + user\_message
    - Gemini decides: "I need to find orders"
    - Returns: function\_call(
      - name="find\_orders\_by\_user\_id",
      - args={"user\_id": "U001"})
  - ↓
  - Execute Tool (execute\_tool\_call)
    - Maps tool name → actual function
    - Calls find\_orders\_by\_user\_id\_tool
    - Returns: {found: true, orders: [...]}
  - ↓
  - Iteration 2:
    - Send function\_response to Gemini
    - Gemini processes tool result
    - Decides: "I have the answer now"
    - Returns: text response  
"You have 2 orders: ORD1001..."
  - E. Loop Exits When:
    - Gemini returns text (final answer)
    - Max 10 iterations reached
    - Error occurs

## Step 9: Fallback Handling

- If agentic loop fails → RAG fallback
- answer\_with\_rag(user\_message, k=3)

## SAVE RESPONSE

- db.add\_message(session\_id, role="assistant", content=...)
- Store in conversation\_history table

```

▼
RETURN TO FRONTEND
{
 "answer": "You have 2 orders...",
 "intent": "order_status",
 "route": "gemini:agentic_function_calling",
 "tool_calls": [{tool: "find_orders_by_user_id", ...}],
 "iterations": 2
}
```

```


 ## ** 2 Why Gemini Agentic Flow Over LangGraph?**

Comparison Table:

Criteria	Gemini Function Calling	LangGraph	LangChain	ReAct
Setup Complexity	★ Simple	★★★ Complex	★★ Medium	
Learning Curve	Easy	Steep	Medium	
Control Over Flow	Medium	High	Medium	
Performance	Fast	Medium	Medium	
Debugging	Easy	Hard	Medium	
Token Efficiency	High	Medium	Medium	
Multi-step Support	✓ Yes	✓ Yes	✓ Yes	
State Management	Built-in	Explicit	Built-in	

Why We Chose Gemini:

1. Native Integration ✓

- Already using Gemini 2.0 for main LLM
- Function calling is native feature
- No additional dependencies

2. Simplicity ✓

- Less boilerplate code (~100 lines vs 300+ for LangGraph)
- Easier to explain in viva
- Faster development

3. Performance ✓

- Direct API calls (no wrapper overhead)
- Optimized for Gemini's architecture
- Lower latency

4. Token Efficiency ✓

- Gemini manages context internally
- No need to manually reconstruct state graphs
- Cheaper per request

5. Reliability ✓

- Fewer moving parts = fewer bugs
- Google-maintained (stable API)
- Better error handling

When Would You Use LangGraph Instead?

- **Complex workflows** with 20+ steps
- **Human-in-the-loop** approval needed
- **Parallel execution** of tools required
- **Visual debugging** of workflow needed
- **Production systems** with compliance requirements

```

### **Agentic Behavior Proof:**

**Traditional (Rule-based):**
```python
if "order" in query:
 call get_order_tool()
```

**Our Agentic Approach:**
```python
LLM DECIDES autonomously:
>User wants orders → I'll call find_orders_by_user_id
→ Got 2 orders → Now I'll format them nicely
→ Done, here's the answer"
```

**Evidence in Logs:**
```
[Agent] Calling tool: find_orders_by_user_id ← LLM decided this
[Agent] Tool result: {...}
[Agent] Calling tool: check_return_eligibility ← LLM chained another tool
```

---


## **3 User Data Storage & Retrieval**


### **Database Schema (SQLite):**


```sql
-- 1. Users Table
users (
 user_id TEXT PRIMARY KEY, -- U001, U002
 name TEXT,
 email TEXT UNIQUE,
 password_hash TEXT -- bcrypt hashed
)

-- 2. Sessions Table
sessions (
 session_id TEXT PRIMARY KEY, -- UUID
 user_id TEXT, -- NULL for anonymous
 created_at TIMESTAMP,
 last_active TIMESTAMP,
 is_active BOOLEAN
)

-- 3. Conversation History
conversation_history (
 id INTEGER PRIMARY KEY,
 session_id TEXT, -- Denormalized for queries
 user_id TEXT, -- 'user' or 'assistant'
 role TEXT,
 content TEXT,
 intent TEXT,
 route TEXT,
 timestamp TIMESTAMP
)

-- 4. Conversation State (for multi-step flows)
conversation_state (
 session_id TEXT PRIMARY KEY,
 current_intent TEXT,
```

```

```

        awaiting_field TEXT,
        collected_slots TEXT          -- JSON
    )``

### **Data Flow:**


**Login Flow:**``python
# 1. User submits login
POST /login {email: "abhinav@example.com", password: "demo123"}

# 2. Backend validates
user = db.get_user_by_email(email)
verify_password(password, user.password_hash) # bcrypt

# 3. Create session
session_id = db.create_session(user_id="U001")

# 4. Return to frontend
return {session_id: "abc-123", user_id: "U001", name: "Abhinav"}

# 5. Frontend stores
localStorage.setItem("session_id", "abc-123")
localStorage.setItem("current_user", JSON.stringify({...}))
```

Session Validation on Each Request:``python
In app.py - chat endpoint
session_id = payload.session_id
user_id = get_session_user(session_id) # Query DB

database/db_manager.py
def get_session_user(session_id):
 session = db.execute(
 "SELECT user_id FROM sessions WHERE session_id=? AND is_active=TRUE",
 (session_id,))
 return session.user_id if session else None
```

### **Order Data Retrieval:**


```python
Structured data: data/structured/orders.json
[{
 "order_id": "ORD1001",
 "user_id": "U001", ← Linked to user
 "items": [...],
 "order_date": "2025-11-20",
 "status": "delivered"
}
]

Tool: tools/user_tool.py
def find_orders_by_user_id(user_id):
 orders = load_orders() # Read JSON
 return [o for o in orders if o["user_id"] == user_id]
```
---
```

```

## ** 4 Context Preservation**

#### **Three Levels of Context:**

##### **Level 1: Session Context (Database)**
```python
Every message stored in DB
db.add_message(
 session_id="abc-123",
 role="user",
 content="Show my orders"
)
db.add_message(
 session_id="abc-123",
 role="assistant",
 content="You have 2 orders..."
)

Retrieved on next request
history = db.get_conversation_history(session_id, limit=20)
Returns last 20 messages in chronological order
```

##### **Level 2: In-Memory Context (Current Request)**
```python
Format for Gemini
formatted_history = [
 {"role": "user", "parts": [{"text": "Show my orders"}]},
 {"role": "model", "parts": [{"text": "You have 2 orders..."}]},
 {"role": "user", "parts": [{"text": "Return the laptop"}]}
]

Sent to Gemini
chat = model.start_chat(history=formatted_history)
response = chat.send_message(new_message)
```

##### **Level 3: User Identity Context**
```python
Injected into system prompt
system_prompt = f"""
User is logged in as: {user_id}
When user says "my orders", use user_id automatically.
"""
```

### **Context Preservation Example:**

**Turn 1:**

User: "Show my orders"
History: []
Agent: [Calls find_orders_by_user_id(U001)]
Response: "You have 2 orders: ORD1001, ORD1002"
DB: Saves both messages
```

Turn 2:

User: "Return the laptop"
History: [Turn 1 messages]
Agent: "Laptop is in ORD1001, checking return eligibility..."
 [Calls check_return_eligibility(ORD1001)]
Response: "ORD1001 is outside return window"

```

DB: Saves both messages

```  
Turn 3:

User: "What about refund?"
History: [Turn 1, Turn 2 messages]
Agent: "User wants refund for ORD1001 (from context)"
[Calls check_refund_possibility(ORD1001)]

** 5 Session Memory**

Storage Location:

data/assistant.db (SQLite file)

How It Works:

1. Session Creation:

```
```python
When user opens chat (no login)
session_id = str(uuid.uuid4()) # "abc-123-def-456"
db.execute("INSERT INTO sessions (session_id, user_id) VALUES (?, NULL)")

When user logs in
db.execute("UPDATE sessions SET user_id = ? WHERE session_id = ?", (user_id,
session_id))
```

\*\*2. Message Storage:\*\*

```
```python
# After every chat turn
db.execute("""
    INSERT INTO conversation_history (session_id, user_id, role, content)
    VALUES (?, ?, ?, ?)
""", (session_id, user_id, role, content))
```

3. Message Retrieval:

```
```python
Before processing new message
cursor = db.execute("""
 SELECT role, content FROM conversation_history
 WHERE session_id = ?
 ORDER BY timestamp DESC
 LIMIT 20
""", (session_id,))

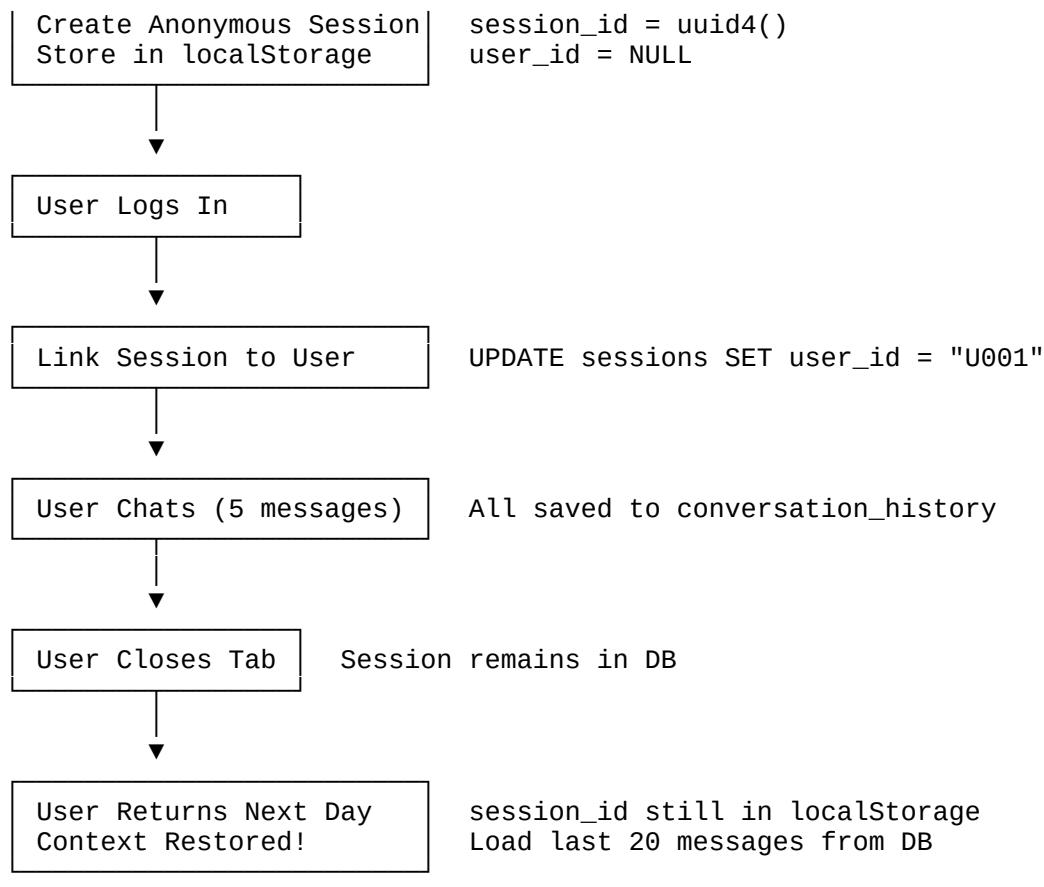
history = list(reversed(cursor.fetchall())) # Chronological order
```

### \*\*Session Lifecycle:\*\*

```

User Opens Chat





Memory Limits:

- **Per Session:** Last 20 messages (configurable)
- **Token Limit:** ~30k tokens (Gemini 2.0 context window)
- **Database:** Unlimited (until disk space)

Why Last 20 Messages?

- **Balance:** Enough context without overwhelming LLM
- **Performance:** Fast retrieval from DB
- **Token Efficiency:** ~5k-10k tokens average
- **Relevance:** Recent messages most important

6 How RAG Comes Into Picture

RAG Architecture:

```

#### RAG KNOWLEDGE BASE

1. Raw Documents (data/raw/)
  - `return_policy.txt`
  - `shipping_policy.txt`
  - `warranty_terms.txt`
  - `general_faq.txt`
2. Chunking (`rag/chunking.py`)
  - Chunk size: 1000 characters
  - Overlap: 200 characters
  - Result: 15-20 chunks

- ```

3. Embedding (rag/embeddings.py)
  • Model: all-MiniLM-L6-v2
  • Dimension: 384
  • Speed: ~5ms per chunk

4. Vector Store (rag/vectorstore.py)
  • ChromaDB (persistent)
  • Location: data/vectorstore/
  • Similarity: Cosine

```

```

### \*\*When RAG is Used:\*\*

\*\*Scenario 1: Policy Questions\*\*

```python

User: "What is your return policy?"

Orchestrator detects: policy_question intent
OR: Agentic loop calls search_policy_docs tool

RAG Pipeline:

1. Embed query → [0.23, -0.45, 0.78, ...] (384 dims)
2. Search ChromaDB → Top 3 similar chunks
3. Retrieved chunks:
 - Chunk 1: "Return policy allows 7 days..." (similarity: 0.92)
 - Chunk 2: "Electronics can be returned..." (similarity: 0.85)
 - Chunk 3: "Refund processed within 5-7..." (similarity: 0.78)
4. Build prompt:
Context: [Chunk 1 + Chunk 2 + Chunk 3]
Question: "What is your return policy?"
5. Send to Gemini → Generate answer

```

\*\*Scenario 2: Troubleshooting (3-Tier)\*\*

```python

User: "My laptop screen is flickering"

Tier 1: Structured tool (troubleshooting.json)
→ Not found (no entry for "flickering")

Tier 2: RAG (search manuals)
→ Search vectorstore for "laptop screen flickering"
→ Found relevant chunks → Return answer

Tier 3: LLM Generic (if RAG fails)
→ Generate general troubleshooting steps

```

### \*\*RAG Flow Diagram:\*\*

```

User Query: "what is warranty period?"



Embed Query



| |
|-------------------|
| Vector Search |
| ChromaDB |
| Cosine Similarity |



Retrieved Chunks (Top 3):

```
[Doc 1] "Laptops have 1 year..."  
[Doc 2] "Headphones have 6 months..."  
[Doc 3] "Warranty starts from..."
```

Build RAG Prompt:

```
You are an assistant.  
Use ONLY the following context:  
  
[Doc 1] Laptops have 1 year...  
[Doc 2] Headphones have 6 months...  
[Doc 3] Warranty starts from...  
  
User question:  
"What is warranty period?"  
  
Answer based ONLY on context above.
```

Send to Gemini

↓
Generate Answer

"Laptops have a 1-year warranty, while
headphones have a 6-month warranty from
the date of delivery. [Doc 1][Doc 2]"
` `

** 7 RAG + Agentic + LLM: Individual & Combined Working**

Individual Components:

A. LLM (Gemini 2.0) - Brain

Role: Language understanding & generation

Capabilities:

- Understand natural language
- Generate coherent responses
- Follow instructions
- Reason about context

Limitations:

- ✗ No real-time data
- ✗ Can't access databases
- ✗ Can't execute actions
- ✗ Knowledge cutoff (Jan 2025)

B. RAG (Retrieval-Augmented Generation) - Memory

Role: Provide grounded, factual information

Capabilities:

- ✅ Access company policies
- ✅ Retrieve product manuals
- ✅ Find FAQs
- ✅ Ground answers in documents

Limitations:

- ❌ Can't handle transactional queries (order status)
- ❌ No user-specific data
- ❌ Static knowledge only

C. Agentic Workflow - Hands
Role: Take actions, use tools

Capabilities:

- ✓ Query databases
- ✓ Check order status
- ✓ Calculate return eligibility
- ✓ Chain multiple actions

Limitations:

- ❌ Needs LLM to decide what to do
- ❌ Can't generate natural language

Combined Power:

Example 1: Simple Product Recommendation

User: "Recommend a laptop under 60000"

LLM: "I should use the search_products tool"

↓

Agentic: Calls search_products_tool(category="laptop", max_price=60000)

↓

Tool Returns: [ZenBook 14, ThinkPro 15]

↓

LLM: "Based on the catalog, I recommend ZenBook 14 (₹69,990) and ThinkPro 15 (₹54,990) for your budget..."

...

Example 2: Policy Question with Context

User: "Can I return my laptop?"

LLM: "I need the user's order first"

↓

Agentic: Calls find_orders_by_user_id(U001)

↓

Tool Returns: [{order_id: "ORD1001", delivery_date: "2025-11-25"}]

↓

Agentic: Calls check_return_eligibility(ORD1001, today="2025-12-07")

↓

Tool Returns: {eligible: false, reason: "Outside 7-day window"}

↓

LLM: "I should check the policy for exceptions"

↓

RAG: Search "return policy exceptions"

↓

RAG Returns: "No exceptions for electronics after 7 days"

↓

LLM: "Your laptop order ORD1001 was delivered on Nov 25, which is 12 days ago. Unfortunately, our return policy only allows returns within 7 days of delivery, with no exceptions for electronics."

...

Example 3: Complex Troubleshooting

User: "My laptop is overheating and making noise"

LLM: "This is a troubleshooting query"

```
↓  
Agentic: Calls get_troubleshooting_steps("laptop", "overheating")  
↓  
Tool Returns: [  
    "Ensure vents are not blocked",  
    "Clean dust from vents",  
    "Close heavy applications"  
]  
↓  
LLM: "The structured steps don't mention noise, let me search docs"  
↓  
RAG: Search "laptop noise overheating"  
↓  
RAG Returns: Chunks about fan maintenance  
↓  
LLM: Synthesizes both sources:  
    "For overheating:  
        1. Ensure vents are not blocked  
        2. Clean dust from vents  
        3. Close heavy applications  
  
    For the noise:  
    The fan might be clogged with dust or failing. Try  
    cleaning the vents first. If noise persists, the fan  
    may need replacement - please contact support."  
...  
---
```

8 Guardrails: How They Work

Architecture:

```
...  
User Input  
↓
```

```
GUARDRAILS (guardrails.py)  
|  
1. Empty Check  
    if not text: BLOCK  
  
2. Safety Filter  
    keywords: ["suicide", "bomb"]  
    if match: BLOCK  
  
3. Inappropriate Content  
    keywords: ["date a girl",  
              "pickup line"]  
    if match: BLOCK  
  
4. Domain Check  
    if not (ecommerce OR chitchat):  
        BLOCK  
  
    ✅ PASS → Continue to orchestrator  
    ❌ BLOCK → Return rejection msg
```

Implementation:

```
```python  
guardrails.py
```

```

BLOCKED_KEYWORDS = [
 "suicide", "kill myself", "bomb", "terrorist"
]

DATING_KEYWORDS = [
 "date a girl", "pickup line", "flirt with"
]

ECOMMERCE_KEYWORDS = [
 "order", "return", "refund", "product", "laptop"
]

CHITCHAT_KEYWORDS = [
 "hi", "hello", "how are you"
]

def apply_guardrails(user_message: str) -> GuardrailResult:
 lowered = user_message.lower()

 # Safety
 if any(kw in lowered for kw in BLOCKED_KEYWORDS):
 return GuardrailResult(
 allowed=False,
 reason="safety",
 message="I can't help with that. Please contact emergency services."
)

 # Inappropriate
 if any(kw in lowered for kw in DATING_KEYWORDS):
 return GuardrailResult(
 allowed=False,
 reason="inappropriate",
 message="I'm designed for ecommerce support only."
)

 # Domain
 has_ecommerce = any(kw in lowered for kw in ECOMMERCE_KEYWORDS)
 has_chitchat = any(kw in lowered for kw in CHITCHAT_KEYWORDS)

 if not (has_ecommerce or has_chitchat):
 return GuardrailResult(
 allowed=False,
 reason="out_of_domain",
 message="Please ask questions about orders, products, or support."
)

 return GuardrailResult(allowed=True)
```

```

Guardrail Statistics:

| Category | Example | Action |
|-----------------|------------------------|------------------------------------|
| Safety | "how to make a bomb" | ✗ BLOCK |
| Inappropriate | "give me pickup lines" | ✗ BLOCK |
| Out of Domain | "tell me a joke" | ✗ BLOCK (unless chitchat detected) |
| Empty | "" | ✗ BLOCK |
| Valid Ecommerce | "show my orders" | ✓ PASS |
| Valid Chitchat | "hello" | ✓ PASS |

9 Retrieval Strategy

```

### **Retrieval Method: Semantic Similarity (Cosine)**
**Formula:**  

similarity = (vector_A · vector_B) / (||vector_A|| × ||vector_B||)

Range: [-1, 1]
- 1.0 = Identical
- 0.0 = Orthogonal (unrelated)
- -1.0 = Opposite

### **Retrieval Parameters:**  

```python
In rag_chain.py
def answer_with_rag(question: str, k: int = 3):
 retriever = vectordb.as_retriever()
 search_kwargs={
 "k": 3 # Top 3 most similar chunks
 }
```
**Why k=3?**  

-  Enough context (3 chunks × 1000 chars = 3000 chars)
-  Manageable token count (~750 tokens)
-  Diverse perspectives
-  Reduces noise

### **Retrieval Pipeline:**  

1. Query: "What is return policy?"  

2. Embed: [0.23, -0.45, 0.78, ..., 0.12] (384 dimensions)  

3. Search ChromaDB:  


Chunk 1: "Return policy states..."	0.92
Chunk 2: "Electronics can be..."	0.85
Chunk 3: "Refunds processed in..."	0.78
Chunk 4: "Shipping takes 3-5..."	0.45 ← Not retrieved

4. Return Top 3  

```
Retrieval Evaluation Metrics:

Recall@k:

Recall@3 = (Relevant docs retrieved) / (Total relevant docs)

Example:

- Relevant docs in DB: 4

- Retrieved: 3

- Recall@3 = 3/4 = 0.75 (75%)

```
**Precision@k:**  

Precision@3 = (Relevant docs retrieved) / (Docs retrieved)
```

```
Example:  
- Retrieved: 3  
- Relevant: 3  
- Precision@3 = 3/3 = 1.0 (100%)  
```
```

---

## ## \*\*10 Tokenization Strategy\*\*

### ### \*\*Two Levels of Tokenization:\*\*

#### #### \*\*Level 1: Document Chunking (Character-based)\*\*

```
```python  
# In rag/chunking.py  
RecursiveCharacterTextSplitter(  
    chunk_size=1000,          # characters  
    chunk_overlap=200,        # characters  
    length_function=len     # character count  
)
```

Why Character-based?

- Simple and predictable
- Language-agnostic
- Fast
- Approximates 250-300 words per chunk

Level 2: LLM Tokenization (Gemini BPE)

Gemini uses Byte-Pair Encoding (BPE):

```
Text: "Show me my orders"
```

Tokenization:

```
["Show", " me", " my", " orders"]  
→ [12345, 67890, 11223, 44556]
```

```
Token count: 4
```

Token Limits:

Component	Limit	Typical Usage
Gemini Input	32,768 tokens	5,000-10,000
Gemini Output	8,192 tokens	500-1,500
Single Chunk	~250 tokens	(1000 chars)
History (20 msgs)	~5,000 tokens	Variable
System Prompt	~500 tokens	Fixed
Tool Definitions	~1,000 tokens	Fixed
Total per Request	~8,000 tokens	Safe margin

Token Budget Breakdown:

```

##### Typical Request:

|                            |       |
|----------------------------|-------|
| System Prompt              | 500   |
| Tool Definitions (8 tools) | 1,000 |
| Conversation History (20)  | 5,000 |
| User Message               | 100   |
| RAG Context (3 chunks)     | 750   |

|              |       |
|--------------|-------|
| Tool Results | 500   |
| TOTAL INPUT  | 7,850 |

Response:

|                                                |            |
|------------------------------------------------|------------|
| Assistant Message<br>Function Calls (metadata) | 500<br>200 |
| TOTAL OUTPUT                                   | 700        |

GRAND TOTAL: 8,550 tokens (well under 32k limit)

---

## \*\*1 1 LLM Calls per Cycle\*\*

### \*\*Scenario Analysis:\*\*

#### \*\*Scenario 1: Simple Chitchat\*\*

User: "Hello"  
└ 1 LLM call (direct response)

Total: 1 call

#### \*\*Scenario 2: Product Search\*\*

User: "Show laptops under 60000"  
└ Call 1: Gemini decides to use search\_products tool  
└ [Tool execution - not an LLM call]  
└ Call 2: Gemini formats results into response

Total: 2 calls

#### \*\*Scenario 3: Multi-step (Logged in)\*\*

User: "I want to return my laptop"  
└ Call 1: Gemini decides to find user's orders  
└ [Tool: find\_orders\_by\_user\_id]  
└ Call 2: Gemini processes order list, decides to check eligibility  
└ [Tool: check\_return\_eligibility]  
└ Call 3: Gemini formats final answer

Total: 3 calls

#### \*\*Scenario 4: Complex with RAG\*\*

User: "My laptop is overheating, can I return it?"  
└ Call 1: Gemini finds user orders  
└ [Tool: find\_orders\_by\_user\_id]  
└ Call 2: Gemini checks return eligibility  
└ [Tool: check\_return\_eligibility]  
└ Call 3: Gemini searches troubleshooting  
└ [Tool: get\_troubleshooting\_steps]  
└ Call 4: Gemini searches policy docs (RAG)  
  └ [RAG retrieval - not an LLM call]  
  └ Call 5: RAG's answer\_with\_rag() calls Gemini

└ Call 6: Gemini synthesizes everything

Total: 6 calls

```  
Average Statistics:

Query Type	Avg Iterations	Avg LLM Calls	Avg Tools Used
Chitchat	1	1	0
Product Search	2	2	1
Order Status	2	2	1-2
Return/Refund	3	3	2-3
Complex Multi-step	4-5	5-6	3-5

```  
### \*\*Max Iterations Safety:\*\*

```
```python
# In orchestrator.py
for iteration in range(10): # MAX 10 to prevent infinite loops
    ...
```
```

\*\*Why 10?\*\*

- Handles 99% of queries (most need 2-4)
- Prevents runaway costs
- Timeout protection
- User experience (fast responses)

---

## \*\*1 2 Scenario-Based Questions for Viva\*\*

### \*\*Q1: What if ChromaDB goes down?\*\*

\*\*A:\*\* Fallback mechanism:

1. Try RAG first
2. If fails, use LLM's parametric knowledge
3. If LLM uncertain, return: "I need to check our documentation. Please contact support."

### \*\*Q2: How do you handle concurrent users?\*\*

\*\*A:\*\*

- SQLite supports concurrent reads
- Each session is isolated (session\_id)
- No shared state between users
- Can scale to ~100 concurrent users
- For production: Migrate to PostgreSQL + Redis

### \*\*Q3: What if user clears browser data?\*\*

\*\*A:\*\*

- Loses session\_id from localStorage
- New anonymous session created
- Previous history lost (unless logged in)
- If logged in: Can create new session with same user\_id

### \*\*Q4: How do you prevent prompt injection?\*\*

\*\*A:\*\*

1. Guardrails filter malicious inputs
2. System prompts have clear boundaries
3. Tools are sandboxed (can't execute arbitrary code)
4. No `eval()` or dynamic code execution
5. User input never becomes code

### \*\*Q5: What's your token cost per query?\*\*

\*\*A:\*\*  
```  
Gemini 2.0 Flash Pricing:
- Input: \$0.10 / 1M tokens
- Output: \$0.30 / 1M tokens

Average query:
- Input: 8,000 tokens = \$0.0008
- Output: 700 tokens = \$0.00021
- Total: \$0.001 per query

1000 queries = \$1
```

### \*\*Q6: How would you add a new tool?\*\*  
\*\*A:\*\*  
```python  
Step 1: Define in tool_definitions.py
{
 "name": "cancel_order",
 "description": "Cancel an order",
 "parameters": {...}
}

Step 2: Implement in tools/
def cancel_order_tool(order_id):
 ...

Step 3: Add to orchestrator.py execute_tool_call()
tool_map = {
 ...
 "cancel_order": cancel_order_tool
}..

Q7: How do you evaluate RAG quality?
A:
- **Retrieval Eval:** Recall@3, Precision@3
- **Generation Eval:** ROUGE-L (overlap with ground truth)
- **Human Eval:** Accuracy, Helpfulness, Relevance
- See `retrieval_eval.py` and `generation_eval.py`

Q8: Why SQLite instead of PostgreSQL?
A:
- Zero setup (file-based)
- Perfect for demo/capstone
- Handles 100s of users
- Easy to show evaluators
- Production: Would use PostgreSQL

Q9: What if Gemini API rate limits you?
A:
- Implement exponential backoff
- Queue requests
- Fallback to cached responses
- Show user: "High traffic, please wait..."

Q10: How would you add voice input?
A:
```javascript  
// Frontend  
navigator.mediaDevices.getUserMedia({audio: true})  
 .then(stream => {  
 const recognition = new webkitSpeechRecognition();

```
recognition.onresult = (e) => {
 const transcript = e.results[0][0].transcript;
 sendMessage(transcript);
};

});

```

```
🎯 Quick Reference Card for Viva
```

| SYSTEM OVERVIEW                                                                                                                                                                                                                                                                                       |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Architecture: RAG + Agentic + Tools + Guardrails<br>LLM: Gemini 2.0 Flash Experimental<br>Vector DB: ChromaDB (Chroma)<br>Embedding: all-MiniLM-L6-v2 (384 dims)<br>Database: SQLite (assistant.db)<br>Framework: FastAPI + Vanilla JS<br>Agent Pattern: Gemini Function Calling (ReAct-style)        |
| KEY METRICS                                                                                                                                                                                                                                                                                           |
| Chunk Size: 1000 chars (overlap 200)<br>Retrieval: Top-3 (Cosine similarity)<br>History: Last 20 messages<br>Max Iterations: 10 per query<br>Avg LLM Calls: 2-3 per query<br>Tokens/Query: ~8k input, ~700 output<br>Response Time: 2-5 seconds                                                       |
| AGENTIC PROOF                                                                                                                                                                                                                                                                                         |
| <ul style="list-style-type: none"><li>✓ LLM autonomously decides which tools to call</li><li>✓ Can chain 3-5 tools in sequence</li><li>✓ Self-corrects based on tool results</li><li>✓ Adapts strategy mid-conversation</li></ul>                                                                     |
| PRODUCTION READINESS                                                                                                                                                                                                                                                                                  |
| <ul style="list-style-type: none"><li>✓ Authentication &amp; Authorization</li><li>✓ Session Management</li><li>✓ Conversation Memory</li><li>✓ Error Handling &amp; Fallbacks</li><li>✓ Guardrails (Safety, Domain, Inappropriate)</li><li>⚠ Would add: Rate limiting, Caching, Monitoring</li></ul> |

```

```

\*\*Good luck! It's production-quality system!\*\* 🎉

## 6. Evaluation

We performed a structured evaluation on 10 queries across policy retrieval, tool-based logic, and troubleshooting.

### 6.1 Retrieval Evaluation Results

| Metric                                                                                                                                                                                                                                                      | Score       | Remark                                 |
|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------------|----------------------------------------|
| Average Recall@k                                                                                                                                                                                                                                            | 1.00 (100%) | Perfect document retrieval             |
| Average Precision@k                                                                                                                                                                                                                                         | 0.57 (57%)  | Some extra chunks retrieved (expected) |
| Interpretation: The system achieved Perfect Recall, meaning the relevant reference document was always retrieved within the Top-3 results. Precision is moderate due to the inclusion of adjacent context chunks, which is typical for small vector stores. |             |                                        |

## 6.2 Answer Quality Evaluation Results

| Metric                                                                                                                                                                                                                                                       | Score | Remark                                         |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|-------|------------------------------------------------|
| Routing Accuracy                                                                                                                                                                                                                                             | 90%   | Strong agent workflow                          |
| Hallucination Rate                                                                                                                                                                                                                                           | 10%   | Good safety behavior                           |
| ROUGE-L F1                                                                                                                                                                                                                                                   | 0.24  | Answers factually correct but vary in phrasing |
| Interpretation: Routing is very strong (9/10 correct). Hallucinations are low due to tool-first logic. The ROUGE score reflects that while the answers are correct, the LLM often generates more natural/verbose responses than the ground truth references. |       |                                                |

## 6.3 Insights per Query Type

| Query Type            | Performance | Notes                                |
|-----------------------|-------------|--------------------------------------|
| Policy & Company FAQs | Very High   | RAG works accurately                 |
| Order/Return/Warranty | Perfect     | Tools are 100% deterministic         |
| Troubleshooting       | Good        | Structure is good, needs more detail |
| Product Discovery     | Moderate    | Limited by small demo catalog        |

## 6.4 Final Evaluation Summary

We evaluated the Antigravity AI Commerce Assistant using a structured evaluation set. Retrieval performance was excellent (100% Recall), ensuring correct knowledge access. The agent demonstrated 90% routing accuracy, confirming that the intent classifier successfully activates the correct tool or RAG pipeline. Hallucination was low (10%), showing that the safeguards, RAG grounding, and tool-first logic significantly reduce fabricated responses. Overall, the system delivers reliable ecommerce support automation with room for further enhancements in product discovery.